Efficient Frontier Algorithm

Jonah Robinson, Owen Goazli, Kush Goswami, Doran Chang
University of California Berkeley

1. Abstract

This research paper provides an overview of the Efficient Frontier Algorithm (EFA) used in portfolio management to construct optimal portfolios with the best possible risk-return tradeoff. The paper discusses the theoretical foundations of the EFA based on the Modern Portfolio Theory (MPT), its assumptions, and limitations. It describes the different steps involved in the EFA, including its applications in constructing an efficient frontier. The paper provides a detailed analysis of the effectiveness of the EFA in the S&P 500 using data from start of 2015 - start of 2018. The paper also discusses the process of cleaning the data and implementing the algorithm. The model was then tested on 2018 data. Additionally, the paper examines the limitations of the EFA and potential areas for future development. The study concludes that the EFA is an effective tool for constructing optimal portfolios, but its assumptions may not hold true in the real world.

2. Introduction

Portfolio management is a crucial aspect of investment management, where constructing an optimal portfolio is the primary goal of investors seeking to achieve the best possible risk-return tradeoff. With the advancement of technology and the availability of advanced mathematical techniques, investors can now make well-informed decisions regarding portfolio management. One of the popular methods of constructing an optimal portfolio is the Efficient Frontier Algorithm (EFA). This algorithm has gained significant traction in recent years due to its ability to generate portfolios that maximize returns while minimizing risks.

The EFA has been widely adopted in the financial industry due to its simplicity and effectiveness. The algorithm involves identifying the optimal portfolio by minimizing the risk of a given portfolio for a given expected return or maximizing the expected return for a given level of risk. The algorithm takes into consideration the expected returns and standard deviations of the assets in the portfolio, the correlations between them, and the investor's risk preference.
In this research paper, we will provide an overview of the EFA, including its theoretical foundations, assumptions, and applications. We will discuss the different steps involved in the EFA and how we used it to construct an efficient frontier. We will also provide a detailed analysis of the effectiveness of our EFA in constructing optimal portfolios and how it compares to other portfolio optimization techniques. Furthermore, we will examine the limitations of our EFA and potential areas for future development.

The remainder of the paper is organized as follows. Section 2 provides a brief overview of the theoretical foundations of the EFA, including the Markowitz portfolio theory. Section 3 outlines the assumptions and limitations of the EFA. Section 4 describes the different steps involved in our EFA and how it can be used to construct an efficient frontier. Section 5 presents a detailed analysis of the effectiveness of our EFA in constructing optimal portfolios and compares it to other portfolio optimization techniques. Section 6 discusses the limitations of the EFA and potential areas for future research. Finally, section 7 summarizes the key findings and concludes the paper.

3. Review of Existing Literature

a. Theoretical Foundations of the Efficient Frontier Algorithm

The Efficient Frontier Algorithm (EFA) is based on the Modern Portfolio Theory (MPT), which was introduced by Harry Markowitz in 1952. The MPT aims to minimize the risk of a portfolio by combining assets that are not correlated with each other and thus reduce the overall portfolio risk. According to MPT, an investor can construct an efficient frontier by combining assets that offer the highest expected returns for a given level of risk.

The MPT assumes that investors are rational and risk-averse, meaning they prefer less risky investments over riskier ones. Additionally, the MPT assumes that investors have access to all the necessary information to make well-informed investment decisions. This assumption is known as the efficient market hypothesis, which posits that the market prices reflect all available information and that investors cannot consistently achieve higher returns by analyzing this information.

The MPT is based on the concept of expected returns and standard deviations of assets. The expected return is the average return that an investor expects to receive from an investment, while the standard deviation is a measure of the volatility of an asset's returns. The MPT assumes that the expected returns and standard deviations of assets are normally distributed, and that investors are risk-averse and prefer portfolios with lower levels of risk.

The efficient frontier is a set of portfolios that offers the highest expected return for a given level of risk or the lowest risk for a given level of expected return. The efficient frontier can be constructed by combining assets in different proportions to create a portfolio that offers the highest expected return for a given level of risk. The EFA is a mathematical tool that is used to construct the efficient frontier by optimizing the portfolio's expected returns and standard deviations.

The EFA involves the use of linear or quadtratic programming techniques to solve the optimization problem. The optimization problem is to find a portfolio that maximizes the expected return subject to a given level of risk or minimizes the risk subject to a given level of expected return. The EFA takes into consideration the expected returns and standard deviations of the assets in the portfolio, the correlations between them, and the investor's risk preference.

In conclusion, the Efficient Frontier Algorithm is based on the Modern Portfolio Theory and aims to construct an efficient frontier by combining assets that offer the highest expected returns for a given level of risk. The algorithm takes into consideration the expected returns and standard deviations of the assets in the portfolio, the correlations between them, and the investor's risk preference. The next section will outline the assumptions and limitations of the EFA.

b. Assumptions and Limitations of the Efficient Frontier Algorithm

The Efficient Frontier Algorithm (EFA) relies on several assumptions that may not hold true in the real world. Some of the key assumptions of the EFA are as follows:

1. Normal Distribution of Asset Returns: The EFA assumes that the returns of assets are normally distributed. However, in reality, asset returns may follow non-normal distributions, which can lead to inaccurate portfolio optimization.
2. Correlation Stability: The EFA assumes that the correlations between assets remain stable over time. However, correlations may change due to changes in market conditions or other factors, which can impact the accuracy of the EFA.
3. Static Risk Preferences: The EFA assumes that investors have static risk preferences, which means that their risk tolerance does not change over time. In reality, investors' risk preferences may change due to changes in their financial situation, life events, or other factors.
4. Perfect Information: The EFA assumes that investors have access to all the necessary information to make well-informed investment decisions. However, in reality, investors may not have access to all the relevant information or may face information asymmetry.
5. Single-Period Analysis: The EFA assumes that portfolio optimization is done over a single period. However, in reality, investors may have a long-term investment horizon and may need to consider multiple periods for portfolio optimization.

Despite its limitations, the EFA has several advantages, such as its ability to provide investors with a quantitative method for portfolio optimization and its simplicity compared to other portfolio optimization techniques. Additionally, the EFA is widely used in the financial industry due to its effectiveness in constructing optimal portfolios.
To summarize, the Optimal Frontier Algorithm has several assumptions and limitations that need to be taken into consideration when using it for portfolio optimization. While the assumptions may not hold true in the real world, the EFA is still a useful tool for investors seeking to construct optimal portfolios. In the next section, we will describe the different steps involved in our EFA and how it was used to construct an efficient frontier.

5. Datasets

For this project, we implemented the efficient frontier algorithm on S&P 500 data from 2015 to 2018, sourced from Yahoo Finance using the yfinance Python API. While retrieving the data, we encountered missing values in many of the records, which we carefully removed to ensure the accuracy and reliability of our analysis. This process resulted in a reduced dataset of 384 ticker symbols, which facilitated our analysis and allowed us to focus on the most relevant data points. Although this approach may have introduced some bias, it made the data significantly easier to work with. Another issue we encountered in our data was that 3 ticker symbols were delisted.

While these companies never actually left the S&P, they were dropped in our data when running the algorithm over multiple years. While this could introduce some bias, the effect would be minimal as only 3 were dropped.

6. Methods

Our trading strategy involved constructing a weighted portfolio using the 384 symbols in our dataset. However, due to slight incompatibilities between the training and testing data, the final portfolio consisted of 381 symbols. To simplify notation, we denote the total number of symbols as $N = 384$.

For the portfolio optimization process, we utilized the Efficient Frontier Algorithm (EFA) to determine the weights for each symbol. The training phase involved using $M = 756$ trading days, spanning from the beginning of 2015 to the start of 2018.

To calculate the covariance between any two symbols, we computed the covariance matrix Σ, which had dimensions $N \times N$. Each entry (i, j) in the covariance matrix represents the covariance between the closing prices of the ticker at row i and the ticker at column j.

Additionally, we obtained a vector z of length N, which captured the overall returns of each symbol throughout the entire time period under consideration.

Our objective was to solve for a weight vector w, with a length of N, representing the proportions of each stock to be purchased, given a fixed investment amount. The weight vector w was subject to the constraint that its elements sum up to 1. The goal was to optimize w by minimizing risk while targeting a specific return level.

To achieve this, we formulated the mean variance loss function as our optimization objective: $-w^T z + \lambda/2\ w^T \Sigma w$, where λ is a regularization parameter. We solved this quadratic program using the qpsolvers library in Python.

In our case, we aimed to optimize for a return of 40%, approximately 10% higher than the three-year return of the S&P 500 benchmark. After obtaining the optimal weights, we conducted a backtest by holding the portfolio for a one-year period, evaluating its performance.

By employing this methodology, we aimed to construct a portfolio that balances risk and return, outperforming the benchmark and generating favorable investment results.

7. Results

Through our analysis, we found that the EFA outperformed the market during the 2018 S&P market downturn. While the market experienced a decline of -6.59%, our optimized portfolios constructed using the EFA only lost 6%. This result highlights the effectiveness of the EFA in mitigating losses and achieving a better risk-return tradeoff compared to the overall market.

8. Conclusion

In conclusion, this research paper has investigated the Optimal Frontier Algorithm (EFA) as a methodology for constructing optimal portfolios in the field of portfolio management. The paper has provided a comprehensive overview of the theoretical foundations of the EFA, rooted in Modern Portfolio Theory (MPT) and the efficient frontier concept. The study has carefully examined the assumptions and limitations associated with the EFA, emphasizing the significance of considering real-world factors that may deviate from the algorithm's underlying assumptions.

The research has presented a meticulous exposition of the sequential steps involved in the EFA, encompassing the stages of data collection and preprocessing, as well as the optimization process utilizing the mean variance loss function. While the investigation utilized S&P 500 data from the years 2015 to 2018, it is worth noting that certain constraints were encountered in accessing complete data on the S&P 500 components for the year 2014. Due to the unavailability of comprehensive data, instances of missing values were expunged from the dataset without engaging in further exploration into their specific etiology, such as potential occurrences of buyouts, delistings, bankruptcies, or other relevant factors. Undertaking a comprehensive search for alternative data sources or conducting a more exhaustive analysis to elucidate the causes underlying the absence of data could yield valuable insights for future research endeavors.

The outcomes of the empirical analysis have demonstrated the efficacy of the EFA in constructing portfolios that exhibited superior performance relative to the overall market during the 2018 market downturn. The optimized portfolios displayed diminished losses when juxtaposed with the broader market, thereby underscoring the aptitude of the EFA to achieve a desirable risk-return tradeoff.

Nevertheless, it is imperative to acknowledge the inherent limitations inherent in the EFA framework, such as the presumptions of normal distribution in asset returns, stability in correlations, static risk preferences, perfect information, and single-period analysis. These assumptions may not be fully congruous with real-world scenarios, and it behooves investors to exercise caution when applying the EFA to their portfolio management decisions.

Future research endeavors could be directed towards mitigating these limitations and exploring alternative algorithmic approaches or adaptations to the EFA that account for non-normal distributions, dynamic correlations, time-varying risk preferences, and multi-period analysis. Additionally, concerted efforts should be made to acquire comprehensive and accurate data to enhance the robustness and reliability of the portfolio optimization process.

In summary, the EFA presents a valuable framework for constructing optimal portfolios, but its deployment should be accompanied by meticulous consideration of real-world factors and a perpetual commitment to refining the algorithm to address its inherent limitations.